

Chapter 4

New evidence for the authorship of case statements and *dicta* in

Gratian's *Decretum*

Paul Evans

20 July 2016

Since the discovery by Winroth in 1996 of a first recension of Gratian's *Decretum*, there has been scholarly controversy over whether the compiler of the first recension was the same person as the compiler of the second recension. Techniques for authorship attribution developed by researchers in the field of computation linguistics make it possible to approach this problem in a new way. Using the Stylometry with R package, I have analyzed those texts from the *Decretum* traditionally attributed to Gratian personally: the case statements and the *dicta*. Principal component analysis (PCA) of the frequencies of function words (prepositions and conjunctions) in the texts suggests that the author of the case statements was not the same person as the authors of either the first- or second-recension *dicta*. PCA also provides evidence concerning the first- and second-recension *dicta* that suggests they are the product of collaborative authorship rather than being the work of either one or two authors.

Twenty years ago, Anders Winroth announced his discovery of the first recension of Gratian's *Decretum* at the Tenth International Congress of Medieval Canon Law in Syracuse, New York. One of the most important questions raised by that discovery is



whether Gratian 1, the compiler of the first recension was the same person as Gratian 2, the compiler of the second recension. It does not appear that the debate over authorship can be settled using currently available evidence. The goal of my project is to find new evidence for the authorship of case statements and *dicta* in Gratian's *Decretum* using computational methods.¹

First I'm going to discuss in general terms the use of stylometry for authorship attribution. Then, I'm going to discuss the methodology and the software that I am using for this project. Finally, I am going to show the results of stylometric analysis of the case statements, the first-recension *dicta*, and the second-recension *dicta* (first without, then with, the *dicta* from *de Penitentia*), and discuss some possible interpretations of those results.

¹ Earlier versions of this chapter were presented as conference papers. "Can Stylometry Provide New Evidence about the Identity of Gratian 1 and Gratian 2?", was presented to the session on Canon Law in the Twelfth and Thirteenth Centuries at the *Rem non novam nec insolitam aggredimur* conference and grand opening of the Stephan Kuttner Institute of Medieval Canon Law at Yale Law School, May 21-22, 2015. "New evidence for the authorship of case statements and *dicta* in Gratian's *Decretum*" was presented to the Classical Sources III session at the Fifteenth International Congress of Medieval Canon Law (ICMCL) at Université Paris II Panthéon-Assas, July 17-23, 2016.

Stylometry is the measurement of style. “Style is a property of texts constituted by an ensemble of formal features which can be observed quantitatively or qualitatively.”²

While style has both qualitative and quantitative aspects, stylometry is concerned only with quantitative aspects of style. One well-established use of stylometry is to attribute authorship. And for the purpose of authorship attribution, the formal linguistic features that stylometry measures are the frequencies of occurrence of function words.

Linguists draw a distinction between function words and content words. Function words are words like prepositions and conjunctions. Content words are words like adjectives, nouns, and verbs. Function words convey meaning by their use in grammatical structure. The Latin conjunction “*sed*” doesn’t mean anything by itself, but rather it places two words or grammatical constructs into an adversative relationship with each other.

Here’s another way of thinking about the distinction: function words are closed-class words and content words are open-class words. Language-speaking communities can and do make up new adjectives and nouns and verbs all the time; so content words are

² J. Berenike Herrmann, Karina van Dalen-Oskam, and Christof Schöch, “Revisiting Style, a Key Concept in Literary Studies,” *Journal of Literary Theory* 9, no. 1 (2015): 44.

an open class that can be added to at will. But new prepositions and conjunctions are almost never added to a language, and they change very slowly over time, if they change at all, and are therefore, for all practical purposes, a closed, finite, class. In theory, that means that you ought to be able to make a list of every function word in a language, although in practice, that's not easy to do.

Evidence from experimental psychology suggests that both authors and readers process function words at an unconscious level.³ The frequency with which a given author uses particular function words is therefore considered to be more or less invariant, making it a reliable authorial signature.

Stylometric analysis of the frequency of function words for the purpose of attributing authorship has had a number of notable successes. The validity of this approach for textual scholarship was firmly established by the work of Frederick Mosteller and David L. Wallace on the *Federalist Papers*. The authorship of 11 of the *Federalist Papers* had been disputed since the early 19th century, with competing claims advanced on

³ Mike Kestemont, "Function Words in Authorship Attribution from Black Magic to Theory?" in *Proceedings of the 3rd Workshop on Computational Linguistics for Literature (CLfL)* (Gothenburg, Sweden: Association for Computational Linguistics, 2014), 59–66.

behalf of Alexander Hamilton and James Madison. In 1944, Douglass Adair, using traditional scholarly methods, settled the dispute largely to the satisfaction of early American historians, determining that Madison was the author of all 11 of the disputed numbers.⁴ In 1964, Mosteller and Wallace confirmed Adair's findings by conducting a stylometric analysis of the frequencies of 70 function words to compare the 11 disputed numbers with numbers securely attributed to Hamilton and Madison.⁵

Let's take a first look at how this kind of stylometric analysis works in practice with some actual data from the first- and second-recension *dicta*, excluding the *dicta* from *de Penitentia*.⁶ "In" is the most frequently occurring word in the *dicta*. There are 1,450

⁴ Douglass Adair, "The Authorship of the Disputed Federalist Papers," *The William and Mary Quarterly* 1, no. 2 (1944): 98–122 and Douglass Adair, "The Authorship of the Disputed Federalist Papers: Part II," *The William and Mary Quarterly* 1, no. 3 (1944): 235–64.

⁵ Frederick Mosteller and David L. Wallace, *Inference and Disputed Authorship: The Federalist*, Addison-Wesley Series in Behavioral Science: Quantitative Methods (Reading, Mass: Addison-Wesley Pub. Co, 1964).

⁶ Including the *dicta* from *de Penitentia* skews the results of this analysis significantly. **Changes come overwhelmingly from first-recension *de Pen. dicta* (9525 words) vs. second-recension *de Pen. dicta* (556 words).**

"In" is the most frequently occurring word in the *dicta*. There are 1,682 occurrences of "in" out of 66,238 words in the first-recension *dicta* (2.54%). There are 431 occurrences of "in" out of 14,811 words in the second-recension *dicta* (2.91%). "In" therefore occurs 14.6% more frequently in the second-recension *dicta* than it does in the first-recension *dicta*.

occurrences of “*in*” out of 56,713 words in the first-recension *dicta* (2.56%). There are 411 occurrences of “*in*” out of 14,255 words in the second-recension *dicta* (2.88%). “*In*” therefore occurs 13% more frequently in the second-recension *dicta* than it does in the first-recension *dicta*.

“*Non*” is the second most frequently occurring word in the *dicta*. There are 1,360 occurrences of “*non*” in the first-recension *dicta* (2.40%). There are 306 occurrences of “*non*” in the second-recension *dicta* (2.15%). “*Non*” therefore occurs 12% more frequently in the first-recension *dicta* than in the second-recension *dicta*.

13% for “*in*” and 12% for “*non*” are significant variations for such common words—it’s not like we’re talking about low-frequency words where a small difference in the count can make for a big difference in percentage.

We could graph the number of occurrences of “*in*” and “*non*” per 100 words of the two samples (from the first- and second-recension *dicta*), with the percentage frequency of

“*Non*” is the second most frequently occurring word in the *dicta*. There are 1,622 occurrences of “*non*” in the first-recension *dicta* (2.45%). There are 314 occurrences of “*non*” in the second-recension *dicta* (2.12%). “*Non*” therefore occurs 15.5% more frequently in the first-recension *dicta* than in the second-recension *dicta*.



“*in*” on the horizontal axis, and the percentage frequency of “*non*” on the vertical axis, and we would have an extremely simplistic visualization of the total variation between the two samples. Now, we are obviously not going to make an attribution of authorship based on the frequencies of only two function words.

Increasing the number of function words one collects data for increases the accuracy of the stylometric analysis (up to a point). The reason that there is a limit to the marginal value of each additional word included in the analysis is that word frequencies in a sample of text fall off approximately according to an observed empirical relationship known as Zipf’s law. If the words in a sample are rank-ordered from most to least frequent, Zipf’s law postulates that, as a first-order approximation, the frequency of each word will be $1/N$ times that of the most frequent word, where N is the rank.⁷

But increasing the number of function words also introduces a new problem. We were able to represent our stylometric analysis of the frequency of “*in*” and “*non*” in the

⁷ Footnote Jockers. Renumber footnotes. I need to provide examples with numbers that agree with the data for *in* and *non* above. In reality, the word frequencies for the *dicta* do not appear to drop off nearly as sharply as Zipf’s law ($1/N$) would predict. I also need to address the issue that as the wordlist extends downward, the mix shifts from disproportionately function words to disproportionately content word. 16 of the first 20 words are function words usable for analysis (even if I discard *si* for other reasons), but I have to go all the way down to 240 to get somewhere in the neighborhood of 50 suitable function words for analysis.

samples from the first- and second-recension *dicta* on a two-dimensional graph. But there will be as many dimensions on the graph as there are function words for which we collect data. And because human beings are not good at visualizing quantitative data in more than three dimensions, we need to find a way to reduce the number of dimensions. This is where the technique of principle component analysis, or PCA, becomes useful.⁸

PCA first combines as many of the raw dimensions as possible into synthetic components on the basis of strong correlations, either positive or negative. For example, going back to the data on the frequencies of “*in*” and “*non*” in the first- and second-recension *dicta*, the two dimensions of the graph could be collapsed into a single component that could be thought of as representing the probability that “*in*” will, and that “*non*” will *not*, occur in a given sample. (And this is, in fact, what the software that I’m using for this project does.)

⁸ For a general introduction to the use of principal component analysis (PCA) in literary stylometric analysis, see Hugh Craig, “Stylistic Analysis and Authorship Studies,” in *A Companion to Digital Humanities*, ed. Susan Schreibman, Raymond George Siemens, and John Unsworth, Blackwell Companions to Literature and Culture 26 (Malden, MA: Blackwell Pub, 2004), 273–88 and Chapter 6 “Style” in Matthew Lee Jockers, *Macroanalysis: Digital Methods and Literary History*, Topics in the Digital Humanities (Urbana: University of Illinois Press, 2013).

Finally, PCA displays the two components that contribute the most to the total variation between the samples, and graphically arranges the samples according to their probability relative to those two components.

An important (and time-consuming) aspect of any project of this nature is corpus preparation. A baseline requirement for carrying out stylometric analysis is the availability of an electronic text. Ideally, I would be working with electronic texts of good critical editions of both the first and second recensions of Gratian's *Decretum*, following consistent orthographic conventions, and encoded in a standard format like TEI P5 XML. The Mellon Foundation-supported project, directed by Anders Winroth, to edit the first recension is making good progress, but is not yet complete enough for me to use on this project. So, I am working with the electronic text of the Friedberg edition that Timothy Reuter and Gabriel Silagi used to produce the *Wortkonkordanz zum Decretum Gratiani* for the MGH.⁹ The MGH e-text is encoded in the obsolete Oxford Concordance Program format.

⁹ Timothy Reuter and Gabriel Silagi, eds., *Wortkonkordanz Zum Decretum Gratiani*, Monumenta Germaniae Historica. Hilfsmittel 10 (München: Monumenta Germaniae Historica, 1990).

Anders Winroth and Lou Burnard of the Oxford Text Archive (OTA) each provided me with copies of the MGH e-text. The copies differed, and I went through an exercise not unlike preparing a critical edition to restore the e-text to the state that Reuter and Silagi intended.

I generated the sample text for the first-recension *dicta* by extracting from the MGH e-text of the Friedberg edition all of the *dicta* listed by Winroth in the appendix of *The Making of Gratian's Decretum*, and by applying the changes to the *dicta* that differed between the first and second recensions.¹⁰ I generated the sample text for the second-recension *dicta* by starting with all the *dicta* in parts 1 and 2 of the Friedberg edition, and then taking away every word that appeared in the first-recension *dicta*. For the case statements, I simply used the text from the vulgate *Decretum* as it appears in the Friedberg edition.¹¹

¹⁰ Anders Winroth, *The Making of Gratian's Decretum* (Cambridge: Cambridge University Press, 2000), 197–227.

¹¹ This is perhaps not entirely satisfactory. It would be more methodologically consistent with the way in which the samples of the first-recension *dicta* were prepared to apply the differences found in Winroth's appendix to the case statements as well, however the differences are quite minimal. The only case statement (*thema*) for which Winroth notes a textual difference is C.19 d.init. (Winroth, 216). The first-recension version of the text omits a 13-word clause added to the second recension version, seemingly for the purpose of piling up descriptive detail. (*unus relictus propria ecclesia eo inuito, alter*

Because stylometric analysis for authorship attribution depends on the frequencies of prepositions and conjunctions, it is important to include enclitics substituting for conjunctions. Every word in the samples with a *-que* ending that is actually an enclitic, and not just part of the word, has been mapped to the word plus the pseudo-conjunction “xque.”¹²

dimissa regulari canonica cenobio se contulit). None of the wordlists used to perform the principle components analyses include any of these 13 words, so the use of the vulgate rather than a proxy first-recension version of the text of C.19 d.init. has no effect on the outcome of these tests.

¹² In the case statements, 1st-, and 2nd-recension *dicta* from Gratian’s *Decretum*, there are 747 occurrences of 79 unique words ending in *-que*. (This does not count 423 occurrences of the word ‘*que*’ itself.) Of those, 498 are occurrences of 19 unique words from Schinke’s 54-word pass list, while 249 occurrences of 60 unique words are not. It is from these 249 words that, according to Schinke, the *-que* ending should be detached as an enclitic.

However, the 249 words include 72 occurrences of 17 unique words ending with the adverbial enclitics *-cumque* or *-cunque*, from which the *-que* ending should not be detached. The 249 words also include a further 149 occurrences of 21 unique false positives:

cumque, *eque* (*aeque*), *namque*, *pleraque*, *plerique*, *plerisque*, *plerumque*, *quinque*, *unamquamque*, *unaqueque*, *unicuique*, *uniuscuiusque*, *unumquemque*, *unusquisque*, *usquequaque*, *utramque*, *utraque*, *utrique*, *utrisque*, *utriusque*, *utrumque*.

This leaves only 28 occurrences of 22 unique words from which the *-que* ending should actually be detached as an enclitic.

False positives over-represent the frequency of occurrence of the ‘*-que*’ enclitic as a conjunction by an order of magnitude. Including all false positives makes ‘xque’ the 37th most frequent word in the sample, while excluding them makes it the 376th most frequent word. There are 55 occurrences of the word ‘*namque*’, the most frequently occurring false positive. Detaching the ‘*-que*’ ending from ‘*namque*’ overstates the frequency of ‘*nam*’, making what is actually the 480th most frequent word appear to be the

Now that we have the preliminaries out of the way, we can take a look at the results. I used the *stylo* R package to generate all of the plots that I’m going to show you today.¹³ R is a statistical programming language.¹⁴ Mike Kestemont, Maciej Eder, and Jan Rybicki of the Computational Stylistics Group developed the package, and Mike Kestemont in particular has been very generous in his technical advice for this project.

130th, while making ‘*namque*’, which is actually the 176th most frequent word in the samples when false positives are excluded, disappear from the list altogether.

¹³ Maciej Eder, Mike Kestemont, and Jan Rybicki, “Stylometry with R: A Suite of Tools,” in *Digital Humanities 2013: Conference Abstracts* (Lincoln, NE: University of Nebraska–Lincoln, 2013), 487–89, <http://dh2013.unl.edu/abstracts/>.

¹⁴ R Core Team, *R: A Language and Environment for Statistical Computing* (Vienna, Austria: R Foundation for Statistical Computing, 2015), <https://www.R-project.org/>.






Figure 1 updated 11 Dec 2015

Here is the plot of a three-way comparison between the case statements, the first-recension *dicta*, and the second-recension *dicta*, excluding the *dicta* from *de Penitentia*.

The case statements are red, the first-recension *dicta* are green, and the second-recension *dicta* are blue. Each of the texts has been divided into 1200-words samples. Principal component 1, along the horizontal axis is 11.2%. Principal component 2, along the vertical axis is 7.3%. That is, PC1 explains 11.2% of the total variation between the samples, and PC2 explains 7.3% of the total variation between the samples. This is good: as a general rule, we want to see a value for PC1 greater than 10% and we want to see a value for PC2 greater than 5%. The most striking feature of this plot is the fact that the case statements are so far away from the *dicta*, and the next step is to take a look at which function words are producing that effect.




Figure 2 updated 11 Dec 2015

Turning on a stylo option called “feature loadings” lets us see how strongly particular words influence the placement of text samples along the PC1 and PC2 axes. The documentation calls this the feature’s “discriminative strength.” For example, we see that “*sed*” and “*non*” are way out on the right of the PC1 axis, while “*unde*” is way down at the bottom of the PC2 axis.

Remember that in our first experiment with counting function words, “*non*”, the second most common word in the samples, was strongly associated with the first-recension *dicta*. Here we see “*non*” on the far right, and in fact the samples from the first-recension *dicta* (but not from the second-recension *dicta*) tend to spread out to the right. Note also that “*in*”, the most common word in the samples, is actually pretty close to the middle. So, it’s not so much that the second-recension *dicta* have more occurrences of “*in*”, it’s that the first-recension *dicta* have fewer.

What is really interesting here is that “*an*” and “*si*” cluster with the case statements, “*an*” very strongly, “*si*” somewhat less so. This makes sense because indirect questions dominate the language of the case statements. It is a question of genre. So the next step in the stylometric analysis is to control for genre by removing the question words “*an*” and “*si*” from the list of function words.






Figure 3 updated 11 Dec 2015

We've now reached the final stage of the three-way comparison between the case statements, the first-recension *dicta*, and the second-recension *dicta*. We are now using the 49 most frequent words on our function list instead of the 51 most frequent words, having commented out "*an*" and "*si*". And even without "*an*" and "*si*", PC1 still explains 10.5% of the total variation between the samples, down slightly from 11.2%. PC2 still explains 7.3% of the total variation between the samples. So, even controlling for genre, the distance between the case statements and the *dicta*—both first- and second-recension—is still quite striking.

To turn to the other interesting aspect of the three-way comparison, you'll note that the second-recension *dicta* in blue cluster strongly to the upper-left quadrant. Now, Mike Witmore, a member of my dissertation committee who isn't an insider with respect to debates about Gratian's *Decretum*, but is very experienced in the use of stylometry with the plays of Shakespeare, was somewhat optimistic on the basis of this evidence that the first- and second-recension *dicta* might be statistically distinguishable.




Figure 4 updated 11 Dec 2015

So, in an attempt to take a closer look at the *dicta* by themselves, I removed the case statements and ran a two-way comparison of 1000-word samples of just the first- and second-recension *dicta*, again, excluding the *dicta* from *de Penitentia*. (Stylo changes the color assignments depending on the number of samples, so in this plot the first-recension *dicta* are red and the second-recension *dicta* are green.) And the results are ambiguous. The PC1 axis is 9%, somewhat under the 10% threshold we would like to see. Also, although we see the second-recension *dicta* clustering mostly to the right of the PC1 axis, the two sets of samples are not separated as cleanly as we'd like to see, and certainly nowhere near as cleanly as the case statements were from the *dicta*.





Figure 5 updated 11 Dec 2015

All of the slides we've seen so far exclude the *dicta* from *de Penitentia*, so before moving on to my conclusion, I do want to quickly show you what the results look like when we include the first-recension *dicta* from *de Pen.* (there are not enough words in the second-recension *dicta* in *de Pen.* to be statistically significant—9,525 vs. 556). Many scholars have observed that *dicta* and canons are poorly separated in *de Pen.* I believe that the unusual dispersion of the samples that you see in this plot is a result of that feature.

Conclusion

Principal component analysis (PCA) of the frequencies of function words (prepositions and conjunctions) in the texts strongly suggests that the author of the case statements was not the same person as the authors of either the first- or second-recension *dicta*. PCA also suggests (less strongly) that the first- and second-recension *dicta* were not the work of either one or two authors, but are more likely to have been the product of collaborative authorship.

On Monday, Anders presented a sketch of what a stemma for the first recension might look like. It suggested that the textual transmission was far more complicated than we may have imagined (or at least may have hoped for). The results I've presented here



today suggest that the question of authorship is potentially as complicated as the question of transmission. I believe that there is enough evidence at least to question assumptions of monolithic authorships (of either the one Gratian or two Gratiens variety). If we cannot satisfactorily answer the question “was there one Gratian or were there two?” it is probably because that is not the right question to ask.

Bibliography

- Adair, Douglass. “The Authorship of the Disputed Federalist Papers.” *The William and Mary Quarterly* 1, no. 2 (1944): 98–122.
- — —. “The Authorship of the Disputed Federalist Papers: Part II.” *The William and Mary Quarterly* 1, no. 3 (1944): 235–64.
- Craig, Hugh. “Stylistic Analysis and Authorship Studies.” In *A Companion to Digital Humanities*, edited by Susan Schreibman, Raymond George Siemens, and John Unsworth, 273–88. Blackwell Companions to Literature and Culture 26. Malden, MA: Blackwell Pub, 2004.
- Eder, Maciej, Mike Kestemont, and Jan Rybicki. “Stylometry with R: A Suite of Tools.” In *Digital Humanities 2013: Conference Abstracts*, 487–89. Lincoln, NE: University of Nebraska–Lincoln, 2013. <http://dh2013.unl.edu/abstracts/>.
- Herrmann, J. Berenike, Karina van Dalen-Oskam, and Christof Schöch. “Revisiting Style, a Key Concept in Literary Studies.” *Journal of Literary Theory* 9, no. 1 (2015): 25–52.
- Jockers, Matthew Lee. *Macroanalysis: Digital Methods and Literary History*. Topics in the Digital Humanities. Urbana: University of Illinois Press, 2013.



- Kestemont, Mike. "Function Words in Authorship Attribution from Black Magic to Theory?" In *Proceedings of the 3rd Workshop on Computational Linguistics for Literature (CLfL)*, 59–66. Gothenburg, Sweden: Association for Computational Linguistics, 2014.
- Mosteller, Frederick, and David L. Wallace. *Inference and Disputed Authorship: The Federalist*. Addison-Wesley Series in Behavioral Science: Quantitative Methods. Reading, Mass: Addison-Wesley Pub. Co, 1964.
- R Core Team. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing, 2015. <https://www.R-project.org/>.
- Reuter, Timothy, and Gabriel Silagi, eds. *Wortkonkordanz Zum Decretum Gratiani*. Monumenta Germaniae Historica. Hilfsmittel 10. München: Monumenta Germaniae Historica, 1990.
- Winroth, Anders. *The Making of Gratian's Decretum*. Cambridge: Cambridge University Press, 2000.